

Project 1: Predicting Metal Ion Concentration from Multi-Parameter Water Samples using KNN Regression and Cross-Validation

Import libraries ¶

```
In [67]: import pandas as pd
from scipy.stats import kendalltau
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import LeaveOneOut
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import somersd
```

Read and visualize the dataset

Note: This dataset differs slightly from the one used in the video lectures.

In this dataset, some mixtures have 3 replicas, while others have 4 replicas.

In the following cell:

- Read the file `water_data.csv`
- Print the dimensions of the dataset (i.e. number of rows and columns) and display the first 5 rows.
- Identify the inputs and the outputs columns.
- Provide the number of mixtures with 3 replicas and 4 replicas, respectively.

```
In [68]: # Here we are reading the dataset
file_path = 'water_data.csv'
df = pd.read_csv(file_path)
```

In [54]: df

Out[54]:

	c_total	Cd	Pb	Mod1	Mod2	Mod3
0	0.0	0.0	0.0	-0.999216	-0.714208	-0.414911
1	0.0	0.0	0.0	-0.990800	-0.714373	-0.238335
2	0.0	0.0	0.0	-0.990539	-0.714125	0.020788
3	14.0	0.0	14.0	-1.001247	-0.713546	0.945465
4	14.0	0.0	14.0	-1.013727	-0.714125	0.569631
...
220	5000.0	4000.0	1000.0	-0.874613	-0.584157	-1.515186
221	5000.0	4000.0	1000.0	-0.875954	-0.604495	-1.516845
222	5000.0	5000.0	0.0	-0.872241	-0.349768	0.499810
223	5000.0	5000.0	0.0	-0.872021	-0.354729	0.777301
224	5000.0	5000.0	0.0	-0.865256	-0.373414	0.712840

225 rows × 6 columns

In [55]: df.head()

Out[55]:

	c_total	Cd	Pb	Mod1	Mod2	Mod3
0	0.0	0.0	0.0	-0.999216	-0.714208	-0.414911
1	0.0	0.0	0.0	-0.990800	-0.714373	-0.238335
2	0.0	0.0	0.0	-0.990539	-0.714125	0.020788
3	14.0	0.0	14.0	-1.001247	-0.713546	0.945465
4	14.0	0.0	14.0	-1.013727	-0.714125	0.569631

In [69]: *# Print dataset dimensions and first 5 rows*
 print(f"Dataset Dimensions: {df.shape}")
 df.head()

Dataset Dimensions: (225, 6)

Out[69]:

	c_total	Cd	Pb	Mod1	Mod2	Mod3
0	0.0	0.0	0.0	-0.999216	-0.714208	-0.414911
1	0.0	0.0	0.0	-0.990800	-0.714373	-0.238335
2	0.0	0.0	0.0	-0.990539	-0.714125	0.020788
3	14.0	0.0	14.0	-1.001247	-0.713546	0.945465
4	14.0	0.0	14.0	-1.013727	-0.714125	0.569631

```
In [71]: # Identifying unique mixtures based on the first three columns (output column)
grouped = df.groupby(['c_total', 'Cd', 'Pb'])

# Counting the number of replicas for each unique mixture
replica_counts = grouped.size().reset_index(name='Replica_Count')

# Merging back with the original dataset to retain the replica count information
df = df.merge(replica_counts, on=['c_total', 'Cd', 'Pb'], how='left')

# Print number of mixtures with 3 and 4 replicas
num_3_replicas = (df['Replica_Count'] == 3).sum()
num_4_replicas = (df['Replica_Count'] == 4).sum()
print(f"Number of mixtures with 3 replicas: {num_3_replicas}")
print(f"Number of mixtures with 4 replicas: {num_4_replicas}")
```

Number of mixtures with 3 replicas: 129
 Number of mixtures with 4 replicas: 96

```
In [58]: # Define input and output variables
X = df[['Mod1', 'Mod2', 'Mod3']].values
y = df[['c_total', 'Cd', 'Pb']].values
```

C-index code

```
In [59]: # In this cell is the function that computes the c-index value based on Somers' D
# Use this function as the evaluation metric in the Leave-One-Out (LOOCV) analysis

def cindex(true, pred):
    s_d = somersd(true, y=pred, alternative='two-sided')
    c_index = (s_d.statistic + 1.0)/2.0
    return c_index
```

Functions

```
In [60]: # In this cell add the functions that you need for the data analysis part.
```

Results for Leave-One-Out cross-validation

```
In [61]: # Here run your script for Leave-One-Out cross-validation and print the correlation
```

```
In [62]: # Initialize Leave-One-Out cross-validation
loo = LeaveOneOut()

# Evaluate C-index for different k values
k_values = [1, 3, 5, 7]
c_index_results = []

for k in k_values:
    # Initialize lists to store true and predicted values for each output
    y_true_all = [[] for _ in range(3)]
    y_pred_all = [[] for _ in range(3)]

    for train_index, test_index in loo.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        # Train Nearest Neighbors Regression model
        model = KNeighborsRegressor(n_neighbors=k)
        model.fit(X_train, y_train)

        # Make predictions
        y_pred = model.predict(X_test)

        # Store true and predicted values for each output
        for i in range(3):
            y_true_all[i].append(y_test[:, i][0]) # Ensure scalar values
            y_pred_all[i].append(y_pred[:, i][0])

        # Compute C-index for each output variable
        c_indices = [cindex(y_true_all[i], y_pred_all[i]) for i in range(3)]
        c_index_results.append([k] + c_indices)

# Convert results to a DataFrame and display
c_index_df = pd.DataFrame(c_index_results, columns=['k', 'C-index c_total',
print(c_index_df)
```

	k	C-index	c_total	C-index Cd	C-index Pb
0	1		0.908283	0.921869	0.880549
1	3		0.914191	0.899591	0.874452
2	5		0.894101	0.861966	0.854261
3	7		0.873729	0.814152	0.835533

Results for Leave-Replicas-Out cross-validation

```
In [63]: # Here run your script for Leave-Replicas-Out cross-validation and print the
```

```

In [64]: c_index_results_lro = []

for k in k_values:
    y_true_all = [[] for _ in range(3)]
    y_pred_all = [[] for _ in range(3)]

    for mixture, group in df.groupby(['c_total', 'Cd', 'Pb']):
        if len(group) > 1: # Ensure there are replicas
            test = group # Use all replicas as test data
            train = df.drop(test.index) # Exclude all replicas from training

            X_train, y_train = train[['Mod1', 'Mod2', 'Mod3']], train[['c_total', 'Cd', 'Pb']]
            X_test, y_test = test[['Mod1', 'Mod2', 'Mod3']], test[['c_total', 'Cd', 'Pb']]

            model = KNeighborsRegressor(n_neighbors=min(k, len(X_train)))
            model.fit(X_train, y_train)
            y_pred = model.predict(X_test)

            for i in range(3):
                y_true_all[i].extend(y_test.iloc[:, i].values)
                y_pred_all[i].extend(y_pred[:, i])

    c_indices = [cindex(y_true_all[i], y_pred_all[i]) for i in range(3)]
    c_index_results_lro.append([k] + c_indices)

# Convert results to a DataFrame and display
c_index_df_lro = pd.DataFrame(c_index_results_lro, columns=['k', 'C-index c_total', 'C-index Cd', 'C-index Pb'])
print("Leave-Replicas-Out Cross-Validation Results:")
print(c_index_df_lro)

```

Leave-Replicas-Out Cross-Validation Results:

	k	C-index c_total	C-index Cd	C-index Pb
0	1	0.818413	0.778427	0.737775
1	3	0.818673	0.761452	0.768948
2	5	0.811702	0.739717	0.747839
3	7	0.816371	0.715309	0.761849

Plot Leave-One-Out and Leave-Replicas-Out Results

Note: You may plot the results as they were presented in the video lecture (refer to MOOC2-Module 2 .pptx slides).

In [72]: *# just to verify the variables*

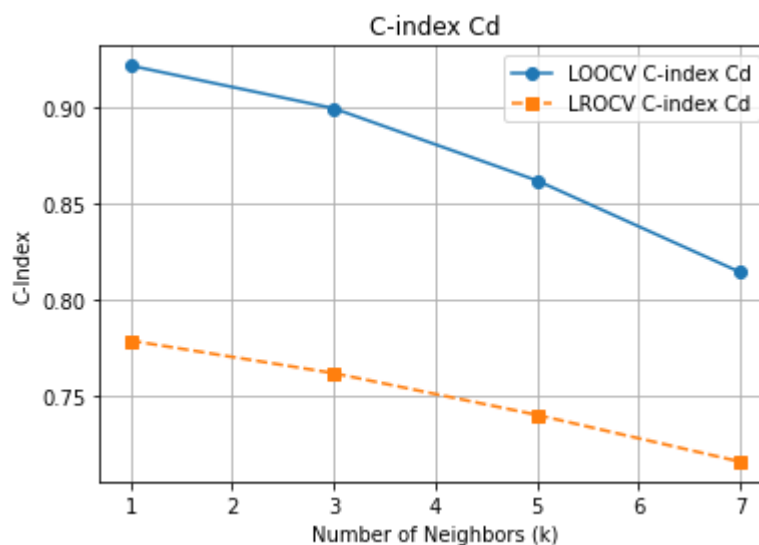
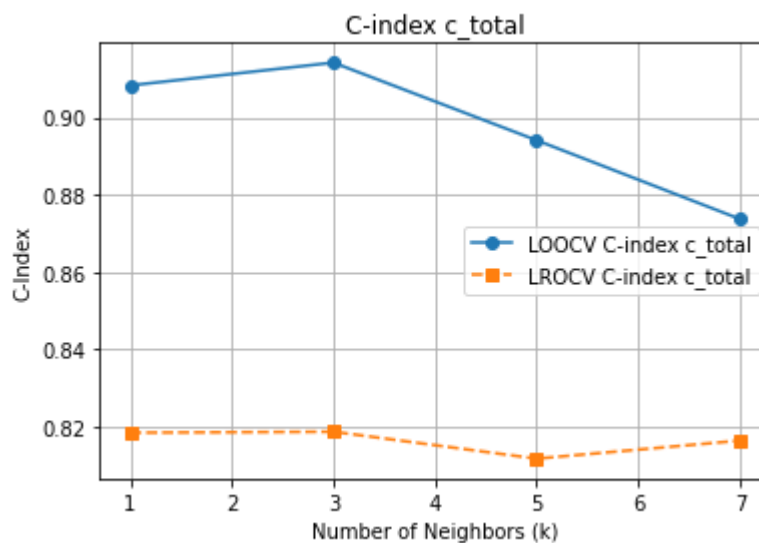
```
print(c_index_df)
print(c_index_df_lro)
```

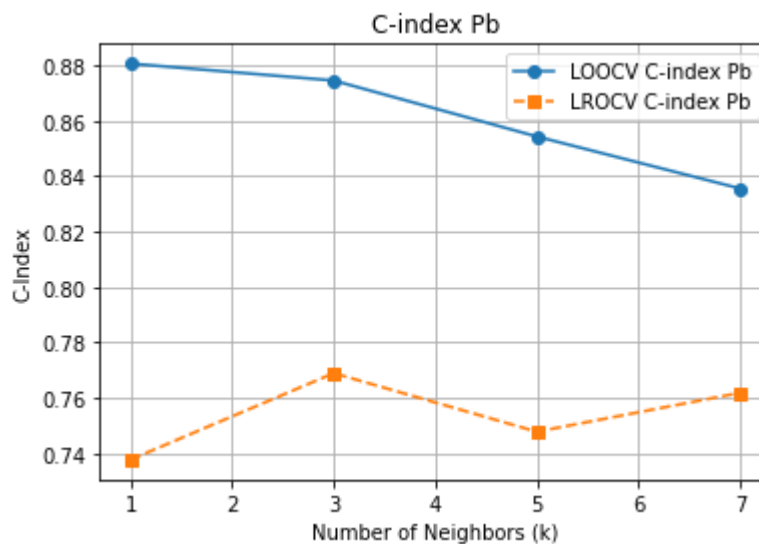
	k	C-index	c_total	C-index Cd	C-index Pb
0	1		0.908283	0.921869	0.880549
1	3		0.914191	0.899591	0.874452
2	5		0.894101	0.861966	0.854261
3	7		0.873729	0.814152	0.835533

	k	C-index	c_total	C-index Cd	C-index Pb
0	1		0.818413	0.778427	0.737775
1	3		0.818673	0.761452	0.768948
2	5		0.811702	0.739717	0.747839
3	7		0.816371	0.715309	0.761849

```
In [66]: #plotting the graphhs between k-valus and c-index for total, Cd and Pb
outputs = ['C-index c_total', 'C-index Cd', 'C-index Pb']
for output in outputs:
    plt.figure(figsize=(6, 4))
    plt.plot(k_values, c_index_df[output], marker='o', linestyle='-', label
    plt.plot(k_values, c_index_df_lro[output], marker='s', linestyle='--',

    plt.xlabel("Number of Neighbors (k)")
    plt.ylabel("C-Index")
    plt.title(f"{output}")
    plt.legend()
    plt.grid()
    plt.show()
```





Interpretation of results

Answer the following questions based on the results obtained

1. Which cross-validation approach produced more optimistic results, and why?
2. Which cross-validation method provides a better estimate of the model's performance on unseen mixtures? Explain your answer.

Answers:

Answer: 1 I think the Leave-One-Out cross-validation (LOOCV) generally produces more optimistic results compared to Leave-Replicas-Out cross-validation (LROCV). This is because LOOCV removes only one single sample at a time, which allows the model to train on almost the entire dataset, which leads to lower variability in predictions and often sometimes higher C-index values. On the other hand, LROCV removes entire sets of replicated measurements for a given mixture, which makes the test set more independent and which in turn prevents the model from relying on nearly identical training samples. This results in a more realistic evaluation of generalization performance, which often leads to lower C-index values compared to LOOCV.

Answer: 2 I think LROCV provides a better estimate of the model's performance on unseen mixtures compared to LeaveLOOCV. I think this is because LROCV removes all replicas of a given mixture from the training set, which ensures that the model is tested on completely unseen data. on the other hand LOOCV only removes one single instance at a time, which means that the test sample often has very similar training samples available, which leads to an overestimation of the model's performance. I think by ensuring that the test data is truly independent of the training set, LROCV is able to give a more realistic measure of the model's generalization ability to new mixtures.